

Distributed Systems - Mandatory 4

Alexander Llewelyn Yan Fatt Soo & Aske Thomas Neye

Github Repo:

<https://github.com/askeneye/ricartagrawala>

System Requirements

R1 (Spec): Implement a system with a set of nodes and a Critical Section that represents a sensitive system operation. Any node may request access to the Critical Section at any time. In this exercise, the Critical Section can be emulated, for example, by a print statement or writing to a shared database on the network

The system consists of three nodes **A**, **B** and **C**, hardcoded on its own port.

A=127.0.0.1:5000, B=127.0.0.1:5001, C=127.0.0.1:5002,

Each node can request to enter the **Critical Section** — represented by a print/log message, which emulates a sensitive operation that must not overlap between nodes.

Each node is a gRPC server (for receiving requests and replies) and a gRPC client (to communicate with peers). Nodes exchange two RPCs:

```
service RAService {
    // A peer asks for permission to enter the critical section.
    rpc RequestCS(Request) returns (Empty);

    // A peer grants permission (may be deferred and sent later).
    rpc SendReply(Reply) returns (Empty);
}
```

R2 (Safety): Only one node may enter the Critical Section at any time

The system satisfies the **safety** requirement with a series of functions. In the **RequestCS()** function, a node defers any incoming request if it is currently executing or has already requested the Critical Section. These deferred requests are only granted after the node exits, as handled by **releaseDeferredReplies()**.

Additionally, a node can only enter the Critical Section with the **TryEnterCriticalSection()** function after receiving replies from *all* other peers, ensuring no overlap occurs.

This combination of deferred replies, Lamport timestamp ordering, and reply counting makes sure that only one node may be in the **Critical Section** at any time.

R3 (Liveliness): Every node that requests access to the Critical Section will eventually gain access

The **Liveliness** requirement is satisfied because the system makes sure that every request for the Critical Section is eventually granted. We handle this with the above-mentioned functions.

In the **RequestCS()** func, requests are never dropped, they are deferred if another node is currently in or waiting for the **Critical Section**. When a node exits, **releaseDeferredReplies()** sends all pending replies, allowing waiting nodes to proceed.

The use of Lamport timestamps, from the file `lamportclock.go`, makes sure we have an ordering of requests, preventing circular waiting or deadlocks. Each node maintains its own **Lamport clock** that increments with every event and message, which ensures global logical ordering of requests.

Discussion:

The system implements the **Ricart Agrawala mutual exclusion algorithm**, using gRPC for passing messages and Lamport timestamps for ordering of requests.

Each node maintains local state (**requesting**, **inCS**, **deferredReplies**) and coordinates access to the **Critical Section** by exchanging **RequestCS** and **SendReply** messages.

The execution logs from Nodes A, B, and C demonstrate that the system satisfies both the **Safety (R2)** and **Liveliness (R3)** requirements.

R2 (Safety):

```
2025/11/12 12:23:23 [A] Requesting critical section (ts=32)
2025/11/12 12:23:23 [C] Sent REPLY(ts=34) to A
2025/11/12 12:23:23 [A] Received REPLY from C (1/2)
2025/11/12 12:23:26 [B] Sent REPLY(ts=35) to A
2025/11/12 12:23:24 [A] Received REPLY from B (2/2)
2025/11/12 12:23:24 [A] ENTERING critical section
2025/11/12 12:23:26 [A] EXITING critical section
```

Because the nodes run on independent local clocks, the wall-clock timestamps in the logs are not perfectly aligned. For instance, as we see Node A's log may show a "Received REPLY" event before the corresponding "Sent REPLY" appears in another node's log. However, the Lamport timestamps (ts=32, ts=34, ts=35) reflect the correct logical ordering of events. These logical clocks, implemented in **LamportClock.Tick()** and **Receive()**, ensure causal consistency across nodes.

R3 (Liveliness):

The logs also show that deferred requests are not lost — they are eventually served once the current holder exits.

For instance, when **C** finishes and releases its deferred replies:

```
2025/11/12 12:23:20 [B] Requesting critical section (ts=29)
2025/11/12 12:23:20 [C] Deferring reply to B
2025/11/12 12:23:22 [C] EXITING critical section
2025/11/12 12:23:23 [C] status: inCS=false, requesting=false, replyCount=2/2, deferred=map[]
2025/11/12 12:23:24 [C] Sent REPLY(ts=32) to B
```

Immediately after, **B** receives this reply and enters its own critical section:

```
2025/11/12 12:23:20 [B] Received REPLY from A (1/2)
2025/11/12 12:23:22 [B] Received REPLY from C (2/2)
2025/11/12 12:23:22 [B] ENTERING critical section
```

This illustrates that deferred replies are released in **releaseDeferredReplies()** after a node exits, ensuring that all waiting nodes eventually proceed.

Appendix – Systems Log

Node A

```
go run ./cmd/node/ -id A -port 5000 -peers A=127.0.0.1:5000,B=127.0.0.1:5001,C=127.0.0.1:5002
2025/11/12 12:22:52 [A] Listening on port 5000
2025/11/12 12:22:52 [A] Server is ready.
2025/11/12 12:22:53 [A] Connected to peer B (127.0.0.1:5001)
2025/11/12 12:22:53 [A] Connected to peer C (127.0.0.1:5002)
2025/11/12 12:22:53 [A] Connected to all peers.
2025/11/12 12:23:01 [A] Requesting critical section (ts=4)
2025/11/12 12:23:01 [A] Received REPLY from C (1/2)
2025/11/12 12:23:01 [A] Sent REPLY(ts=3) to B
2025/11/12 12:23:01 [A] Received REPLY from B (2/2)
2025/11/12 12:23:01 [A] ENTERING critical section
2025/11/12 12:23:03 [A] status: inCS=true, requesting=false, replyCount=2/2, deferred=map[]
2025/11/12 12:23:03 [A] EXITING critical section
2025/11/12 12:23:05 [A] Sent REPLY(ts=10) to C
2025/11/12 12:23:09 [A] Requesting critical section (ts=13)
2025/11/12 12:23:09 [A] Received REPLY from C (1/2)
2025/11/12 12:23:09 [A] Sent REPLY(ts=12) to B
2025/11/12 12:23:09 [A] Received REPLY from B (2/2)
2025/11/12 12:23:09 [A] ENTERING critical section
2025/11/12 12:23:11 [A] Deferring reply to C
2025/11/12 12:23:11 [A] EXITING critical section
2025/11/12 12:23:13 [A] status: inCS=false, requesting=false, replyCount=2/2, deferred=map[]
2025/11/12 12:23:13 [A] Sent REPLY(ts=22) to B
2025/11/12 12:23:13 [A] Sent REPLY(ts=20) to C
2025/11/12 12:23:15 [A] Requesting critical section (ts=23)
2025/11/12 12:23:15 [A] Received REPLY from C (1/2)
2025/11/12 12:23:15 [A] Received REPLY from B (2/2)
2025/11/12 12:23:15 [A] ENTERING critical section
2025/11/12 12:23:17 [A] EXITING critical section
2025/11/12 12:23:20 [A] Sent REPLY(ts=29) to C
2025/11/12 12:23:20 [A] Sent REPLY(ts=31) to B
2025/11/12 12:23:23 [A] status: inCS=false, requesting=false, replyCount=2/2, deferred=map[]
2025/11/12 12:23:23 [A] Requesting critical section (ts=32)
2025/11/12 12:23:23 [A] Received REPLY from C (1/2)
2025/11/12 12:23:24 [A] Received REPLY from B (2/2)
2025/11/12 12:23:24 [A] ENTERING critical section
2025/11/12 12:23:26 [A] EXITING critical section
exit status 0xc000013a
Manual Exit in terminal
```

Node B

```
go run ./cmd/node/ -id B -port 5001 -peers A=127.0.0.1:5000,B=127.0.0.1:5001,C=127.0.0.1:5002
2025/11/12 12:22:52 [B] Listening on port 5001
2025/11/12 12:22:52 [B] Server is ready.
```

```
2025/11/12 12:22:52 [B] Connected to peer A (127.0.0.1:5000)
2025/11/12 12:22:52 [B] Connected to peer C (127.0.0.1:5002)
2025/11/12 12:22:52 [B] Connected to all peers.
2025/11/12 12:22:59 [B] Requesting critical section (ts=1)
2025/11/12 12:22:59 [B] Received REPLY from A (1/2)
2025/11/12 12:22:59 [B] ENTERING critical section
2025/11/12 12:22:59 [B] Received REPLY from C (2/2)
2025/11/12 12:23:01 [B] Deferring reply to A
2025/11/12 12:23:01 [B] EXITING critical section
2025/11/12 12:23:02 [B] status: inCS=false, requesting=false, replyCount=2/2, deferred=map[]
2025/11/12 12:23:03 [B] Sent REPLY(ts=7) to A
2025/11/12 12:23:07 [B] Sent REPLY(ts=9) to C
2025/11/12 12:23:07 [B] Requesting critical section (ts=10)
2025/11/12 12:23:07 [B] Received REPLY from C (1/2)
2025/11/12 12:23:07 [B] Received REPLY from A (2/2)
2025/11/12 12:23:07 [B] ENTERING critical section
2025/11/12 12:23:09 [B] Deferring reply to A
2025/11/12 12:23:09 [B] EXITING critical section
2025/11/12 12:23:11 [B] Sent REPLY(ts=19) to C
2025/11/12 12:23:11 [B] Sent REPLY(ts=17) to A
2025/11/12 12:23:12 [B] status: inCS=false, requesting=false, replyCount=2/2, deferred=map[]
2025/11/12 12:23:13 [B] Requesting critical section (ts=20)
2025/11/12 12:23:13 [B] Received REPLY from A (1/2)
2025/11/12 12:23:13 [B] Received REPLY from C (2/2)
2025/11/12 12:23:13 [B] ENTERING critical section
2025/11/12 12:23:15 [B] Deferring reply to A
2025/11/12 12:23:15 [B] EXITING critical section
2025/11/12 12:23:17 [B] Sent REPLY(ts=26) to A
2025/11/12 12:23:20 [B] Requesting critical section (ts=29)
2025/11/12 12:23:20 [B] Received REPLY from A (1/2)
2025/11/12 12:23:22 [B] Sent REPLY(ts=28) to C
2025/11/12 12:23:22 [B] Received REPLY from C (2/2)
2025/11/12 12:23:22 [B] ENTERING critical section
2025/11/12 12:23:22 [B] status: inCS=true, requesting=false, replyCount=2/2, deferred=map[]
2025/11/12 12:23:23 [B] Deferring reply to A
2025/11/12 12:23:24 [B] EXITING critical section
2025/11/12 12:23:26 [B] Sent REPLY(ts=35) to A
Exit status 0xc000013a
Manual Exit in terminal
```

Node C

```
go run ./cmd/node/ -id C -port 5002 -peers A=127.0.0.1:5000,B=127.0.0.1:5001,C=127.0.0.1:5002
2025/11/12 12:22:52 [C] Listening on port 5002
2025/11/12 12:22:52 [C] Server is ready.
2025/11/12 12:22:52 [C] Connected to peer A (127.0.0.1:5000)
2025/11/12 12:22:53 [C] Connected to peer B (127.0.0.1:5001)
2025/11/12 12:22:53 [C] Connected to all peers.
2025/11/12 12:22:59 [C] Sent REPLY(ts=3) to B
2025/11/12 12:23:01 [C] Sent REPLY(ts=6) to A
2025/11/12 12:23:03 [C] status: inCS=false, requesting=false, replyCount=0/2, deferred=map[]
```

2025/11/12 12:23:05 [C] Requesting critical section (ts=7)
2025/11/12 12:23:05 [C] Received REPLY from B (1/2)
2025/11/12 12:23:05 [C] ENTERING critical section
2025/11/12 12:23:05 [C] Received REPLY from A (2/2)
2025/11/12 12:23:07 [C] EXITING critical section
2025/11/12 12:23:07 [C] Sent REPLY(ts=13) to B
2025/11/12 12:23:09 [C] Sent REPLY(ts=15) to A
2025/11/12 12:23:11 [C] Requesting critical section (ts=16)
2025/11/12 12:23:11 [C] Received REPLY from B (1/2)
2025/11/12 12:23:11 [C] Received REPLY from A (2/2)
2025/11/12 12:23:11 [C] ENTERING critical section
2025/11/12 12:23:13 [C] status: inCS=true, requesting=false, replyCount=2/2, deferred=map[]
2025/11/12 12:23:13 [C] Deferring reply to B
2025/11/12 12:23:13 [C] EXITING critical section
2025/11/12 12:23:15 [C] Sent REPLY(ts=25) to A
2025/11/12 12:23:15 [C] Sent REPLY(ts=23) to B
2025/11/12 12:23:20 [C] Requesting critical section (ts=26)
2025/11/12 12:23:20 [C] Received REPLY from B (1/2)
2025/11/12 12:23:20 [C] ENTERING critical section
2025/11/12 12:23:20 [C] Received REPLY from A (2/2)
2025/11/12 12:23:20 [C] Deferring reply to B
2025/11/12 12:23:22 [C] EXITING critical section
2025/11/12 12:23:23 [C] Sent REPLY(ts=34) to A
2025/11/12 12:23:23 [C] status: inCS=false, requesting=false, replyCount=2/2, deferred=map[]
2025/11/12 12:23:24 [C] Sent REPLY(ts=32) to B
exit status 0xc000013a